

ADL v0.92.2 — Activation Gate Assessment

Not the TAIL-05/#920 external review. See §1.

Verdict

`not_proven`

This matches the position the handoff already states: “*The current truthful verdict is `not_proven`.*” Nothing in this assessment changes it.

1. What this document is, and what it is not

The v0.92.2 handoff states plainly: “*handoff ready; external review has not started*” and “*No external reviewer has been contacted, no disclosure has been authorized, and the external review has not started.*” It defines six activation gates that must **all** be true **before review begins**.

Three of those gates are unsatisfied by the handoff’s own text (§2), and a fourth fails on evidence read from the clone. Performing the external review would therefore breach the packet’s entry conditions. This document assesses the gates instead.

This report is not eligible for intake as the #920 external assessment. The intake contract requires the `adl.external_review_assessment.v1` schema with “*exactly one evidence-backed entry for every one of the 27 accepted #919 finding IDs.*” No finding disposition could be verified here (§4), so no such list exists. Copying this into `findings.json` as a completed external review would manufacture the record the packet exists to prevent.

2. Activation gate assessment

1	Every #919 finding has truthful repair or retained-risk disposition and an independent exact-head re-review	Not satisfied. The handoff itself: “ <i>The 27-finding repair mapping does not itself establish fresh independent execution of all repaired behaviors.</i> ” It further directs: “ <i>Before external dispatch, record each finding’s repair commit, focused proof, independent exact-head re-review and any accepted residual risk.</i> ”
2	#917 and #918 provide accepted handoff, candidate and artifact identities	Partially satisfied / not verifiable here. #918 identities are stated. The retained #917 handoff is partly superseded (podcast deferral, #671/#1169 via PR #1174). No verification possible without Git.
3	Post-remediation candidate revision and artifact-manifest SHA-256 agree across handoff, internal re-review and manifest	Not verifiable, and contradicted locally. See finding P1-A. Clone main is <code>d2f605ce...</code> , not the bound candidate <code>d4575d1b...</code> SHA-256 of the publication manifest could not be computed in this environment.
4	Reviewer independent of implementation and internal-review authorship, with retained evidence	Not satisfied. This lane is operator-invoked, operator-scoped, and operator-consumed. Claude lanes have contributed review artifacts to this project across prior milestones, including material ingested into <code>tail-05/findings.json</code> at v0.92.1. Independence from internal-review authorship cannot be asserted, and no retained independence evidence exists.

5	External contact and exact disclosure scope explicitly authorized	Not satisfied. Handoff: <i>"no disclosure has been authorized."</i>
6	Every required source accessible without exposing material outside approved scope	Not established. Required sources must be read at bound revisions via <code>git show</code> ; this lane has no Git and can read only working-tree bytes.

The handoff's own rule applies: *"If any gate changes, stop and refresh the packet."* Four gates are not in a state that permits review to begin.

3. Findings

P1-A — Supplied working tree is not the bound candidate

Evidence	.git/refs/heads/main = d2f605ce72a8e497c28339271c2bd7a455f30d71. Handoff candidate = d4575d1b34df78b80e71233e5d72cca9bc4339c2. Handoff source revision = 982b6461.... Documentation-audit integration revision = b0896ac1.... The clone matches none of them.
Impact	Any source finding produced from this checkout would describe bytes that are not the candidate's. The handoff anticipates exactly this: <i>"Read packet files with git show <sha>:<path> so a local checkout cannot substitute newer bytes."</i> This lane can only read working-tree bytes, so that protection cannot be honored.
Remediation	Provide a checkout at d4575d1b..., or dispatch a lane with Git execution able to read each packet blob at its bound revision.

P1-B — Review cannot begin: activation gates 1, 4 and 5 are unsatisfied by the packet's own text

Evidence	S2 above, quoting the handoff directly on independent exact-head re-review, disclosure authorization, and reviewer independence.
Impact	Producing an external-review verdict now would record a completed review that the packet states has not started and is not authorized to start. Gate 4 in particular cannot be satisfied by any Claude lane in this configuration.
Remediation	Satisfy gates 1, 4, 5 and 6 and retain the independence and authorization evidence as separate typed records, then dispatch to a reviewer who is not the operator and not an author of prior review artifacts in this repository.

P2-A — Repository split has not occurred; crate count increased

Evidence	Workspace root at the supplied checkout contains adl/, adl-v2/, csdlc-v2/, csdlc-v3/, adl-runtime/, adl-runtime-kernel/, adl-characterization/, adl-resilience/, adl-spec/, plus two additions since v0.92.1: adl-uts/ and adl-provider-core/.
Impact	The multi-generation coexistence raised at v0.91.7, v0.91.8 and v0.92.1 (prior TPR-005, owner #837) persists and has widened. The stated intent to split the repository is not reflected in this candidate.
Remediation	Carry the boot-path map forward as a blocker: for each subsystem the Beta 1 claims rely on, record which crate executes on the candidate's path. Retire superseded generations against parity evidence rather than adding alongside them.

P2-B — Code-reduction claim is again unverifiable, and was not assessed

Evidence	The milestone request asserts the codebase was <i>"reduced considerably."</i> No recursive before/after measurement bound to the candidate was located. The directory-size method this lane used in prior milestones was self-reported as non-proving because relocation into subdirectories registers as deletion; it is not reused here, and no substitute measurement was available for non-candidate bytes.
-----------------	---

Impact	The claim remains unsupported by retained evidence, as at v0.91.8 (prior TPR-004, owner #836).
Remediation	Publish a recursive line count across both revisions with the exact command, paths, and deleted-versus-relocated split, bound to the candidate.

4. Internal finding dispositions

Zero of 27 verified. The #919 register is immutable evidence at report revision 6fc19987... against reviewed source 5c4a6149.... Verification requires reading those blobs at those revisions, which needs Git. The four repair groups (A/#1161, B/#1162, C/#1159, D/#1160) are recorded as merged with named merge commits, and the handoff asserts each merge is an ancestor of the bound candidate — none of which could be checked here. The handoff is explicit that this mapping is not sufficient: *“Issue creation and work in progress do not establish a fix,”* and the repair mapping *“does not itself establish fresh independent execution of all repaired behaviors.”*

5. Verified non-findings

The handoff does not overclaim, and states its own adverse verdict

It records *“The current truthful verdict is not_proven,”* declares the external review not started, distinguishes the frozen internal baseline from the post-remediation candidate, and warns *“Do not silently reinterpret the frozen findings against changed bytes.”* It also holds Group B's disposition map open as *“a source mapping, not a current integration-status record.”*

The quality gate refuses promotion and names specific anti-patterns

QUALITY_GATE_v0.92.2.md records *not proven*, states Beta 1 cannot enter publication finalization without current retained evidence against the integrated candidate, and that *“Partial success is recorded by track and cannot be promoted to a milestone pass.”* The atomic delivery gate rejects *“schemas, scaffolds, undefined slices, outlines and zero-execution qualification,”* and gate 11 states *“CLI-only or fixture-only proof cannot satisfy #915.”*

Freshness invalidation is defined rather than assumed

“A remediation changing the candidate or published claims invalidates affected review/proof/manifests... refresh internal/external review before the release decision.” This is the control that makes P1-A a detectable condition rather than a silent one.

The follow-on is kept outside the frozen denominator

CI-1160-001 and the #1171/#1172 integration-tooling repair are both recorded as separate from the 27 frozen #919 findings, preserving the denominator.

6. Methods and sources accessed

File reads only. **No command executed:** no Git, build, test, validator, provider or cloud operation. No resource or provider use. Sources read at the supplied checkout (d2f605ce..., not the candidate):

ADL_v0.92.2_EXTERNAL_REVIEW_HANDOFF.md;

docs/milestones/v0.92.2/QUALITY_GATE_v0.92.2.md; .git/refs/heads/main; workspace root listing.

7. Limitations, inaccessible evidence, omitted scope

- **No Git.** Every `git show / git cat -file` step the packet requires was impossible. Blob identity, ancestry of the four repair merges, and revision existence are all unverified.

- **No digest computation.** The publication-manifest SHA-256 (6872e631...) and the five repair source-map digests were not computed or compared.
- **Wrong bytes.** All reads are from a non-candidate revision (P1-A).
- **Omitted entirely:** the #919 full-review packet (`final_report.md`, `run_manifest.json`, `findings.json`, `remediation-groups.json`, `artifact-manifest.json`, `acceptance-inventory.json`); `review-manifest.json`; the #920 documentation audit and coverage inventory; `RELEASE_PLAN`; `MILESTONE_CHECKLIST`; the candidate diff; all source code; CodeFriend Beta 1 behavior; security, privacy, provenance and credential boundaries; deployment claims.
- **Unresolved uncertainty:** whether the repairs work. Four groups merged and closed within seconds of their merges is consistent with automation, and says nothing either way about behavioral proof — which the handoff itself declines to claim.

8. Note on the requested format

The 15-part scorecard was not used. The handoff specifies its own eight-part response contract, which is the governing format for this packet; and a numeric score across categories that were not exercised — here, nearly all of them — would assert coverage this lane does not have. The scored template remains appropriate for a review that actually traverses those categories at the bound candidate.